

AWS

AWS Solns. Architect. SAA - CO2: (dummy).

classmate

Date _____

Page _____

Amazon Web Services → Cloud Provider.

> AWS Account created ✓

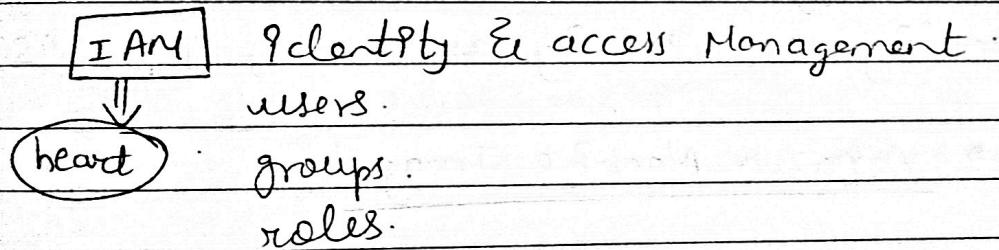
* Module 3:

AWS Regions.

↳ AZ (Availability zones)

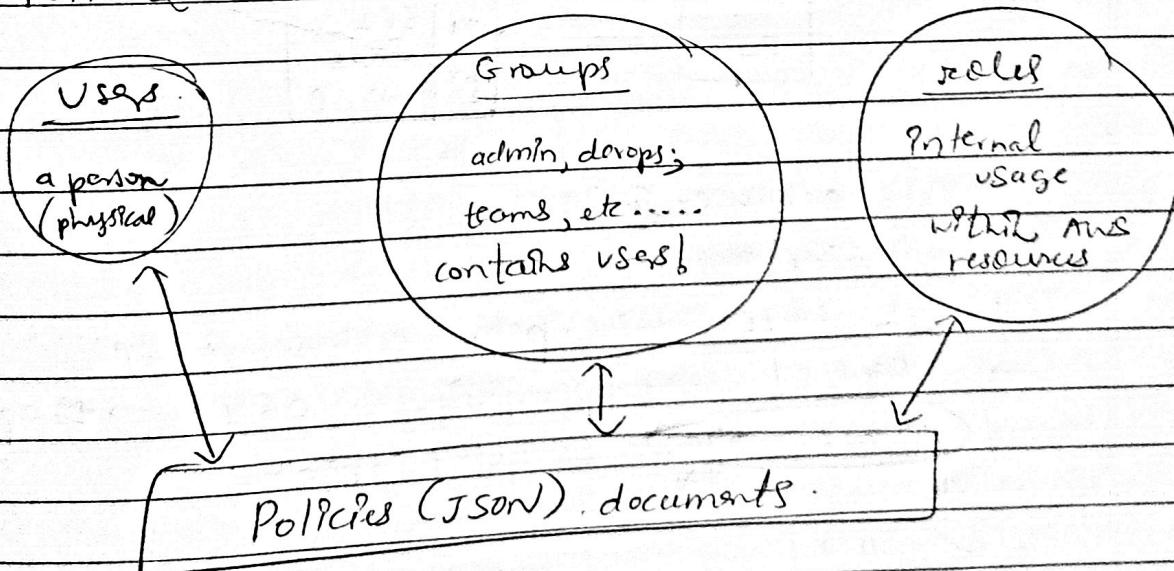
> physical data center.

> all services except IAM & S3.



Root account must never be used (is shared).

> Policies are written in JSON.



> IAM has a global view.

> MFA → multi-factor authentication.

> least privilege principle → to user.

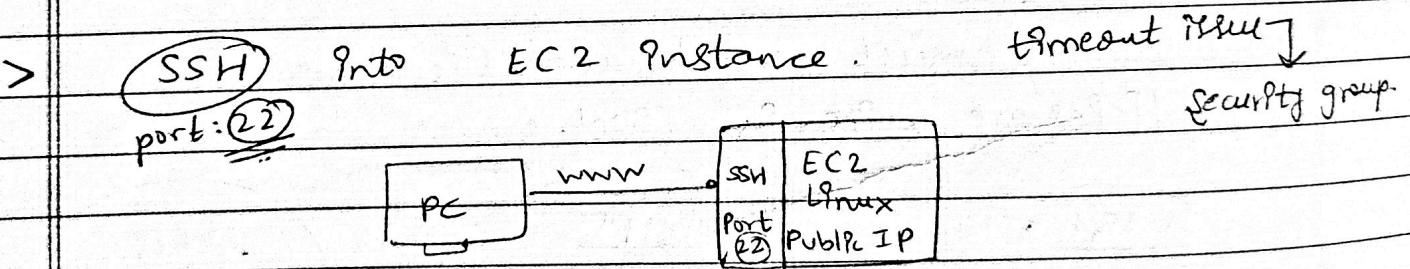
(B9) Enterprise \rightarrow IAM Federation.
 \hookrightarrow SAML standard (active directory).

- > 1 IAM User per person.
- > 1 IAM Role per appn.
- * Never write IAM credentials in code. EVER!
Never use root account.
> all policies \rightarrow 6 tch8 \rightarrow updated ✓

30 Aug. EC2 \rightarrow fundamental! [Elastic Compute Cloud]

- +2 months free
- \rightarrow renting VM's (EC2)
 - \rightarrow storing data on virtual drives (EBS)
 - \rightarrow distributing load across machines (ELB)
 - \rightarrow scaling services using auto-scaling grp (ASG)

(AMI) \rightarrow Amazon Machine Image.



\$ ssh ec2-user @ "IP"

\rightarrow Denied!

\$ ssh -i EC2Tutorial.pem ec2-user @ "IP"

\rightarrow Denied! \Rightarrow permission \Rightarrow 0644 \rightarrow open!

\$ chmod 0400 EC Tutorial.pem

\rightarrow "

✓ Done!

\$ (ctrl+D) or exit

> Security Groups:

↳ controls inbound / outbound traffic.
acts as "firewall" on EC2 instances.
ports, IPv4, IPv6 ✓

- > can be attached to multiple instances.
- > locked down to a region.

> firewall outside EC2 instance.
maintain one separate security group for SSH access.

by default → Inbound traffic → blocked.
outbound traffic → authorized.

> Private vs Public IP (IPv4)

IPv4: 1.160.10.240

IPv6: 3ffe:1900:4545:3:200:f8ff:fe21:67cf
preferred for IoT.

IPv4 ⇒ 3.7 billion addr. in public space.

IPv4: [0-255]. " . " . "

Public: → unique.

→ www access ✓

→ geo located easily.

Pvt. → unique across pvt. network.

→ identified by pvt. network.

→ 2 diff. pvt. networks (2 companies)

can have same IP's. (no problem✓)

→ www access thru Internet gateway (proxy)

→ specified range of IPs (pvt)

Elastic IPs: → to not change IP when on/off EC2.

↳ you own it (until you delete it)

max(5) IPs

↳ 1 instance at a time.

• avoid using them.

- > Instead; use random public IP & register a DNS name to it.
or:
use a load balancer & don't use public IP.

Install Apache on EC2:

→ Login to EC2 ✓

→ sv ✓ (sudo sv)

→ yum update -y ✓

yum install -y httpd.x86_64

systemctl start httpd.service

systemctl enable httpd.service

curl localhost:80

echo "Hello" > /var/www/html/index.html

echo "Hello world from \$(hostname -f)" > /var/www/html/index.html

ip-172-31-43-112.region....

EC2 User Data:

boots trapping EC2 user data.

↳ runs only once at first start.

→ runs with root user.

~~use~~ launch new instance:

user data as text:

#!/bin/bash

||

||

||



EC2 Instance Launch Types:

- > On demand ✓
- > Reserved instances (≥ 1 year). ✓
- > Convertible Reserved instances.
- > Scheduled Reserved Instances.
- > Spot Instances. ✓
- > Dedicated Instances. (no sharing).
- > Dedicated Hosts. (entire physical server) ✓
- Spot fleet

Instance Types: (main ones):

R: RAM needs → In memory caches.

C: CPU needs → compute/dB.

M: balanced ('medium') → general/web app.

I: I/O storage → dB.

G: GPU → video render/ML.

T2/T3: burstable instances (upto capacity) ↴ $\frac{\text{burst}}{\text{compute}}$

T2/T3: unlimted! ✓

to monitor ⇒ cloudwatch.

➤ AMI: machine image. ✓ → (for a specific region).

custom → pre-installed packages

→ fast boot.

→ monitor.

→ security.

→ maintenance.

→ active dir. Integration.

→ using someone else's AMI.

AMI storage → S3 → durable, cheap.

by default AMI is → pvt (can be made publ/pb).
cost → very low.

> Cross Account AMI Copy ✓

- Sharing doesn't affect owner's LIP; but if copied onto new AMI in diff. zone; yes, ownership is affected.
- to copy => Launch EC2 from your acc - using shared AMI to create AMI from instance.

> Placement Groups:

- 1) cluster → 1 AZ, low latency. (same hardware).
- 2) spread → max 7 instances per AZ (hardware spread).
→ for critical apps. (high availability).
- 3) partition → instances spread on diff. partitions
→ which rely on diff. set of racks;
→ scales 100's of EC2 per grp. (Hadoop, Cassandra, Kafka).

ENI → Elastic Network Interfaces

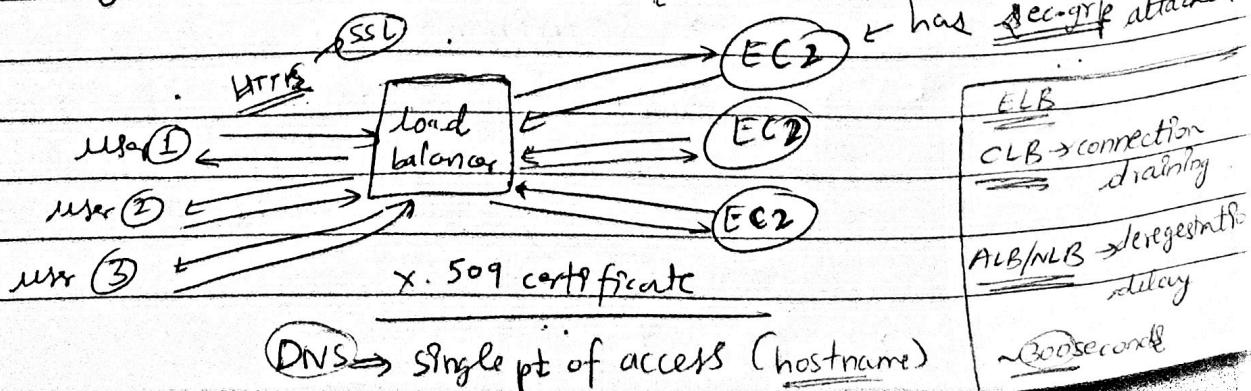
> Scalability:

vertical

horizontal = elasticity → window → high availability.

> Load Balancing: → has static host name.

servers that forward internet traffic to multiple servers (EC2 instances) downstream.



DNS → single pt of access (hostname)

~300 seconds

SSL → Secure Socket Layer
TLS → Transport Layer Security

health checks every 5 sec!

classmate

Date _____
Page _____

- > classic LB - v1. (2009)
- > Appn LB - v2 (2016)
- > Network LB - v2 (2017)

private (priv) or external (public)

- > Appn LB: → microservices & container based appn. (Docker, ECS, ...)
 - ↳ 1 LB many appns ✓ → HTTP traffic
- > supports HTTP/HTTPS & websockets protocol.
- > IP of client not seen directly.

True IP of client → X-Forwarded-For ✓

→ dynamic host port mapping with ECS.

- > Network LB → forwards (TCP) traffic. → SSL termination
 - ↳ less latency ~100ms ⇒ private static IP ↗
 - ↳ sees client IP → higher performance. ↗ static IP per AZ.

→ don't resolve & use underlying IP. cross-zone LB

> 4xx → client induced errors.

5xx → appn. induced errors.

- > LB stickiness: can cause imbalance (50:50 → ideal 80:20 → may happen)

↳ works with Classic LB & Appn. LB.

↳ used for cookies → session data.

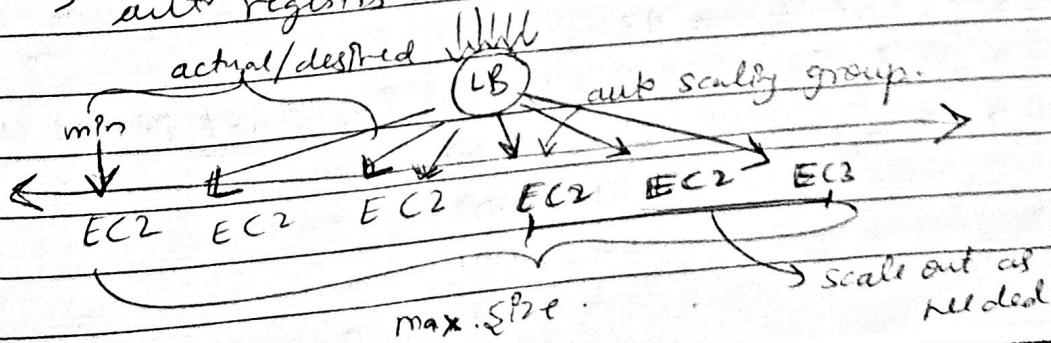
ACM → AWS Certificate Manager

* SNI → "server name indication" → specifies hostname.
SSL/TLS certificates.

> Auto Scaling Group: (ASG).

→ scale in/out EC2 to decrease/increase load.

→ auto registers new instances to LB.



> ASG's: smart hat!

Launch config:

> AMI + instance type

> EC2 user Data.

> EBS volumes.

> Security groups.

> SSH key pair.

min / max / optimal capacity

network / subnet info.

load balancer info.

Scaling policies.

> auto scaling alarms

Cloudwatch

tells when to scale up/down.

> auto scaling rules ✓

* > ASG default termination policy.

(A) + (B) ✗

↳ AZ → most instances ✓ (→ A).

(Delete the one with oldest launch config ✓)

default cooldown → 300 sec.

(✓1) not (✓2) ✗

→ most effective ⇒ 180 sec. ✓

→ Scaling Policies:

lifecycle hook (React? ✗ D.)

(pending state)

(terminating state)

> Target Tracking Scaling:

> Simple / Step Scaling:

> Scheduled Actions:

Launch template vs Launch config

↳ multi-use

↳ re-created each time

> EBS (Elastic Block Storage) → not a physical drive

→ Network drive you can attach to your Instances.

> Locked to an AZ

> billed for provisioned capacity (not PAY AS YOU GO!)

> 4 types: GP2 (SSD)

IOL (SSD)

ST1 (HDD)

SC1 (HDD)

ENI → elastic network interface.

classmate

Date _____

Page _____

> Only GP2 & IO1 → used as boot volumes.

\$ lsblk // show all attached drives

\$ sudo mkfs -t ext4 /dev/xvdb

\$ sudo mkdir /data

\$ sudo mount /dev/xvdb /data.

\$ lsblk

\$ cd /data

\$ sudo ~~cd~~ touch Hello.txt

\$ sudo ls

\$ sudo nano Hello.txt

Hello world!

\$ cat Hello.txt

\$ sudo cp /etc/fstab /etc/fstab.orig

\$ sudo nano /etc/fstab

....
/dev/xvdb /data ext4 defaults,nofail 0 2

\$ cat /etc/fstab

\$ sudo file -s /dev/xvdb

O/P: ext4 ✓

\$ cd ..

\$ sudo umount /data

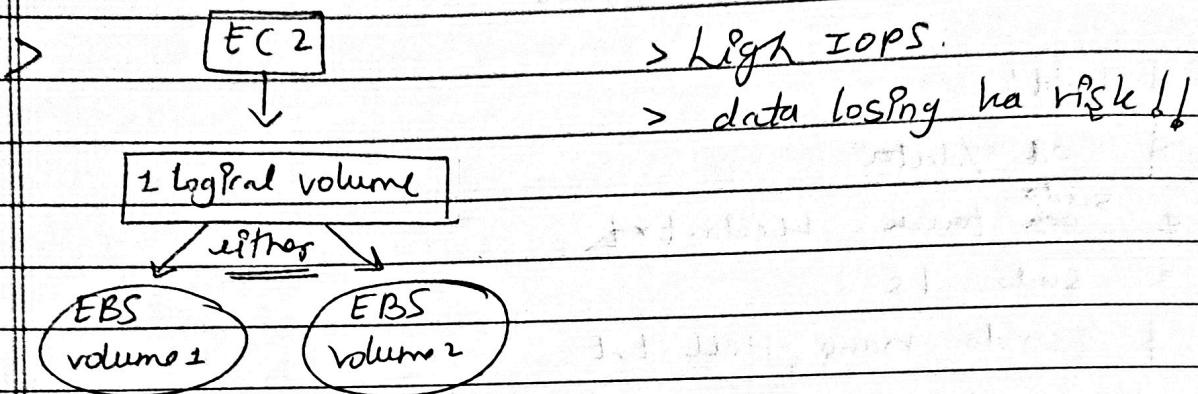
\$ lsblk // drive will be unmounted.

\$ sudo mount -a

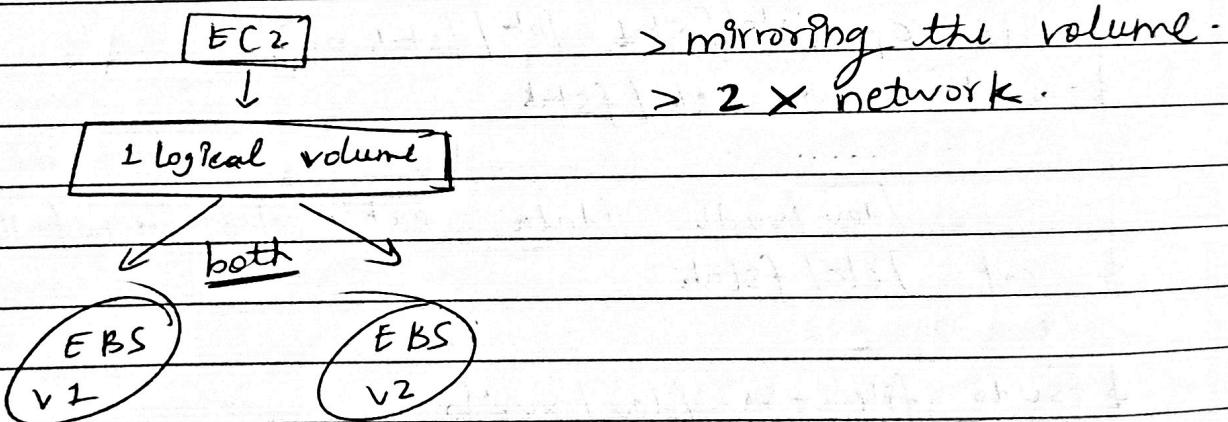
\$ lsblk
// xvdb → disk "/data"

- > EBS snapshots.
- > EBS volume migration.
- > EBS encryption. AES-256 → NMS
- > EBS volume vs Instance store [physical attached disk]
- > EBS RAID options.

RAID 0 (Increase performance)



RAID 1 (Increase fault tolerance)



- > EFS → Elastic File System
- > managed NFS (Network file system)
- > 3x gp2 price (∴ POSIX)
- > pay per use ✓
- > compatible with (Linux) AWS (NOT WINDOWS)
- performance model : 1) general purpose.
- 2) max I/O .

Storage tier : 1) Standard → frequent access
 2) Infrequent access (EFS IA) → cheap

Global Aurora
PostgreSQL → 5432 port
MySQL → 3306 port
↳ RTO → recovery time objective. < 1min

classmate

Date _____

Page _____

> RDS, Aurora, ElastiCache:

Relational Database Service ⇒ SQL ✓

- ↳ PostgreSQL, MySQL, Oracle, MariaDB, Microsoft SQL Server.
- ↳ Managed Service ⇒ automation ✓
- ↳ Storage → EBS backed (IOPS or GP2).
- ↳ It can't (SSH) into instances.

Read Replicas for read scalability.

↳ up to 5 ; async ✓

↳ can be promoted to own DB.

network cost

↳ from 1 AZ to another ; same AZ → free.

RDS multi AZ → elastically recoverable.

↳ SYNC ; 1 DNS name ; more availability; not used for scaling.

Read Replica → multi AZ? ⇒ YES ✓ for disaster recovery.

DB download ⇒ SQL Electron ✓

→ RDS Security → Encryption. ↗ at rest ↗ in-flight

↳ AWS KMS - AES 256

↳ if master not encrypted, read replicas not [can't be replicated].

IAM authentication → for MySQL & PostgreSQL only.

→ Aurora: 5x SQL, 3x PostgreSQL

↳ 20% more cost [RDS]

↳ 6 copies across → 3 AZ. ↗ self-healing & shit.

↳ same as RDS security.

↳ Aurora Serverless? (pay per sec)

→ ElastiCache: → RDS for cache's.

↳ in-memory DB's, high performance, low latency.

↳ makes appn. stateless.

↳ sharding, read replicas ✓

write

read

SSL In-flight ✓

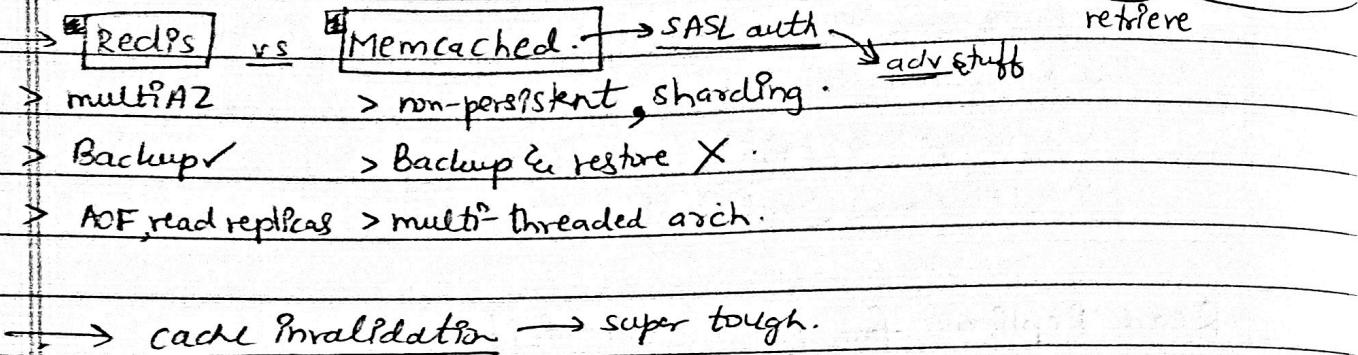
↳ no IAM Auth. but has RDS Auth.

> sharding → multi-node for partitioning of data. classmate
redis → 6379 port

Date _____
Page _____

- DB cache

- User Session Store



> Route 53 → managed DNS (domain name system)

↳ global service ✓

most common records:

> A → hostname to IPv4.

> AAAA → hostname to IPv6

> CNAME → hostname to hostname.

> Alias → hostname to AWS resource.

> \$0.50 per month per hosted zone; ∴ not free.

nslookup xyz.abc.com

→ gives IPv4

// on mac. → dig xyz.com

EC2 meta-data

EC2-AVAIL-ZONE = \$ curl ----- // 169.254.169.254 -----

→ DNS records TTL (time to live)

high TTL → 24 hrs.

low TTL → 60 sec.

→ Root not allowed X

→ CNAME works for 'something.mydomain.com'

↳ 'non-root' domain

→ Alias → xyz.amazonaws.com

↳ Root + Non-Root ✓

→ free + native health check

- Routing Policy → Simple. weighted. Latency. Failover. Geolocn.
- Health Checks → unhealthy if fails 3 checks. • Multivalue.

Route53 as a Registrar → 3rd party. GoDaddy, etc...
 ↳ !DNS x

Sdns. Architect Discussions:

5 pillars ⇒ Cost, Performance, Reliability, Security, Operational Excellence.

Elastic Beanstalk ⇒ developer centric view of deploying appln. on AWS

↳ 3 components:

- ① Appln.
- ② Appln. version: for each deployment.
- ③ Env. name (e.g: dev, test, prod, ...) → free naming.

* Amazon S3 → buckets, directories

↳ global unique name // region level defined
 ↳ consist of objects → have a key → full path.

Key = prefix + object name.

e.g.: s3://my-bucket/my-folder/my-file.txt
 : prefix : obj. name

max. obj. size → 5TB!

version ID ✓

upload limit → 5GB (at a time); metadata; tags ✓

→ S3 versioning. → bucket level.

→ S3 encryption. → SSE-S3, SSE-KMS, SSE-C, client-side.

SSE → server side encryption.

SSEC → HTTPS is mandatory. ✓

→ Encryption in transit (SSL/TLS).

↳ or flight.

→ S3 security: User Based, Resource Based.

↳ ACL → access control list.

IAM policy.

→ pre-signed URLs

add /*

at the end.
classmate

Date _____
Page _____

[ARN] → amazon resource name.

→ JSON based policies ✓

↳ resources: buckets & objects

↳ actions: set of API to Allow or Deny.

↳ effect: Allow / Deny.

↳ principal: acc/user to apply policy to...

→ [CORS] → Cross Origin Resource Sharing.

↳ protocol, domain (host), port.

> AWS CLI, SDK, IAM Roles + Policies:

> Never put credentials on EC2 machine! ✓ bad practice!
↳ use IAM roles instead!

> AWS policy generator, simulator.

EC2 instance metadata: → instances learn about themselves!

URL ⇒ <http://169.254.169.254/latest/meta-data>

AWS SDK > Exponential Backoff ⇒ wait for 2x more than last API call,
if API call fails.

> Advanced S3 & Athena:

MFA-Delete → possible only w/ CLI & not UI.
↳ also only Root?

S3 Access Logs → for auditing (using Athena?)

S3 Replication (CRR, SRR) ⇒ cross region replication,
↳ sync replication? same region replication.

Presigned URLs: to access pvt S3 bucket object for a
specific time to allowed users.

→ S3 storage classes:

S3 Standard - General

S3 Standard - IA (Infreq. access)

S3 Onezone - Infreq. access (IA)

S3 Intell. Tiering

Glacier & Deep Archive.

S3 Object Lock | Glacier Vault Lock
→ worm write once
read many

classmate

Date _____

Page _____

→ server lets!

- > S3 Lifecycle rules → to efficiently switch S3 classes.
> Athena → S3 data analyzer

> CloudFront & AWS Global Accelerator:

↳ CDN (Content delivery network)

distributes cache reads w/ help of edge locations.
→ enhanced security: OAI (origin access identity)

> Signed URL / Cookie

> Global Accelerator

> unicast ; anycast IP's

↳ for low latency (global) ... 1 IP > nearest IP

client gets

> AWS Storage Extras:

Snowball → physical transfer. (Snowball Edge → new).

Snowmobile → Truck | xD

Snowball → S3 → Glacier



> AWS Storage Gateway → S3 access on-premise!

↳ File, Volume, Type Tape

> Amazon FSx (File Server for Windows) + (for Lustre) Linux + cluster.
↳ EFS can't run on Windows (only Linux systems)
↳ (for ML, HPC)

> Decoupling Apps: SQS, SNS, Kinesis, Active MQ:

SQS queue Simple Queue Service

> standard queue.

> msg visibility timeout.

> dead letter queue. ☠

> delay queues.

> FIFO queue egs: Demo fifo • fifo → must!

URN → amazon resource name.SNS: Simple Notification Service:

> Protocol endpoints: → HTTP, HTTPS, Email, Email-JSON, Amazon SQS, Lambda

> SNS + SQS: Fanout | send 1 receive all.

> SNS can't send msg to FIFO queue!

Kinesis → managed alternative to Apache Kafka.

↳ big data 'real time'

→ Streams, Analytics, Firehose.

↳ divided into shards/partitions

> shard-iterator, records, ...

> Amazon MQ → managed message broker service. [Active MQ]

↳ endpoints: Openwire, AMQP, STOMP, MQTT, WS

→ might have SG issue ✓

→ FaaS (fn. as a service).

> Serverless Overview:

↳ don't need to manage servers -

↳ [Lambda, DynamoDB, Cognito, API gateway, S3, SNS & SQS,

Kinesis data Firehose, Serverless Aurora, Step Fn, Fargate.]

AWS Lambda: > virtual functions + auto-scaled.

> short executions + on-demand.

> pay per req & compute time.

Languages: → Node, Python, Java, C#, Go, Ruby, Rust. [DOCKER]

λ limits: memory allocn: 128 MB - 3008 MB

max execution time: 900 sec (15 mins)

env. variables: 4 KB

/tmp size: 512 MB

concurrent exec: 1000

deploym. size: 50 MB

uncompressed size: 250 MB

Lambda@Edge: deploy λ fn alongside CloudFront CDN

OLTP → online transaction processing.
OLAP → " analytical "

classmate

Date _____
Page _____

→ 3 AZ ✓

DynamoDB: Serverless NoSQL dB.

↳ provisioned throughput.

xSplice

WCU → write capacity unit.

RCU → read capacity unit.

DAX → DynamoDB Accelerator → seamless cache.

Streams → CRUD operations.

API Gateway: public REST API's.

↳ easy integration w/ Lambda, HTTP or any AWS service.

Cognito: confusing? → gives users identity to interact w/ app.

↳ User Pools

↳ Identity Pools (Federated Identity).

↳ Cognito Sync (AppSync).

SAM (Serverless Appln. Model)

> Classic serverless REST API: HTTPS, API Gateway, Lambda, DynamoDB

> Databases in AWS:

RDS Aurora ElastiCache DynamoDB S3

Athena Redshift Neptune Elasticsearch

5 pillars of DB: (COSCPR)

Operations Security Cost Reliability Performance

> Monitoring & Audit.

CloudWatch → dashboards → includes graphs.

↳ Log, Agent, Alarms (OK, Insuff. Data, Alarm), Events

CloudTrail → enabled by default. (Record API calls).

Config → audits & records: compliance & config changes.

> IAM Advanced:

> STS: "Security Token Service" → backbone of services!

↳ valid for 1 hr. (grants temporary, limited access).

↳ assumed roles, cross account access.

Identity Federation → 3rd party management

↳ SAML 2.0

Cognito: provide direct access from client side. (mobile, web app).

SAML → security assertion markup lang.

⇒ Directory services: EC2 running Windows?
Microsoft Active Directory (AD).

AWS AD

1) AWS Managed Microsoft AD.

2) AD connector.

3) Simple AD.

opt. 4) Cognito user pools.

⇒ Organizations: global service.

↳ to manage multiple accounts.

SCP (Service Control Policy) → whitelist or blacklist IAM actions

↳ not applied to Root/Master account.

⇒ IAM conditions → if explicit deny exists, it superimposes any underlying allow & vice versa.

⇒ AWS RAM (resource access manager) → avoid resource duplication
↳ share pvt subnet on VPC ✓

⇒ AWS SSO (single sign on) → login once, access all accounts!
↳ SAML 2.0, AD integrated ✓

⇒ Security & Encryption:

KMS (key management service)

→ CMK (Customer Master Key) ⇒ symmetric (AES-256)

→ Asymmetric (RSA & ECC key pair)

SSM Parameter Store: ^{Simple} System Manager Agent

Secrets Manager

CloudHSM > Hardware Security Model

Shield: DDoS prevention

WAF: Web Appln. Firewall (layer 7 → HTTP)

↳ deploy on ALB, API Gateway, CloudFront

Web ACL → SQL inj, XSS prevention?

AWS Shared Responsibility Model

→ cost: use **pvt IP** for lower cost.

Networking VPC:

CIDR - IPv4 :

AWS reserves 5 IP's in each subnet. (1st 4're last addr) → **classless Inter-Domain Routing**

Internet Gateway: connects VPC instances w/ Internet.

Route Table:

NAT Instances: "Network Address Translation" (dell way)

pvt subnet → NAT Gateway → Internet gateway

Network ACL → (inbound/outbound rules → stateless).

→ like firewall (controls traffic from subnet to subnet).

→ can block IP at subnet level.

by default, everything is denied.

VPC Peering: connects 2 VPC's privately.

VPC Endpoint: access services with pvt network.

→ 2 ways: ① Interface ② Gateway.

→ One AWS CLI cmd → region <your-region> (soz default is us-east-1).

VPC Flow Logs + Athena

→ IP traffic info

Bastion Hosts → SSH into pvt. instances.

→ SG must be super strict. (port 22 only).

It is a public subnet (instance).

Site-to-site VPN links VPN Gateway & Customer Gateway.

Direct Connect + DC Gateway

→ dedicated pvt connection. → any IP/any port → all public addrs!

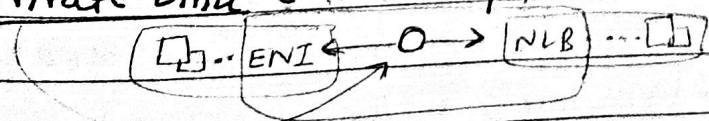
Egress (outgoing) Only Internet Gateway

→ IP/Port (only)

NAT → IPv4 ✓

→ edit route tables too ✓

AWS Private Link (VPC Endpoint Service)



EC2 Classic, AWS ClassicLink (deprecated) → VPC has a ✓.

VPN CloudHub: provide secure commn. bw sites on multi VPC connections.

Transit Gateway: transitive peering bw 1000s of VPC's star connection.

→ supports IP multicast.

(hub & spoke)

NLB → doesn't have SG

classmate

Date _____

Page _____

→ Disaster Recovery:

RPO → Recovery Point Objective.

RTO → Recovery Time Objective.

DMS (Database Migration Service): to migrate dBs ✓

↳ continuous data replication ⇒ CDC (change data capture)

Aws SCT (Schema Conversion Tool): convert database schema from 1 engine to another.

SMS (Server Migration Service): on-premises

VM Import/Export; DNS ✓

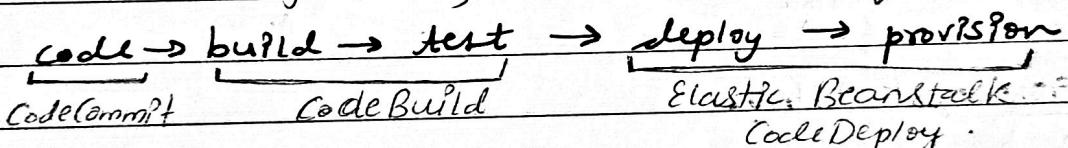
DataSync: move large amount of data from on-premises to AWS.

↳ 1) NFS/SMB ⇒ S3, EFS, FSx 2) EFS ⇒ EFS

→ Other Services:

CI/CD: CodeCommit, CodeBuild, CodeDeploy

↳ continuous integration ; continuous delivery.



⇒ orchestrate: AWS CodePipeline

CloudFormation: Infrastructure as code.

↳ StackSets: CRUD stacks across multiple regions with a single operation!

ECS (Elastic Container Service): run Docker on EC2.

(ECR (Elastic Container Registry))

→ serverless? ⇒ Fargate ✓

EKS (Elastic Kubernetes Service): managed k8s clusters on AWS ⇒ ✓ ; ≈ ECS but diff. API

Step Functions: serverless, fn (multiple) ✓

SWF (Simple Workflow Service) → emit signal, child process

OSCPR

classmate

Date _____

Page _____

EMR (Elastic MapReduce): helps create Hadoop clusters.

→ (Big Data) analysis.

Aws Glue: fully managed ETL (extract, transform, load)

↳ serverless✓, automated code-generation!

Aws Opsworks: managed Chef & Puppet.

↳ manage configs.

Aws Elastic Transcoder: convert media files (video + music).

↳ jobs, pipeline, presets, notifications.

↳ serverless✓.

(virtual desktop
infrastr.)

Aws Workspaces: managed secure Cloud Desktop. (VDI)✓

Aws AppSync: store & sync data across mobile + web apps.

↳ real-time✓ ⇒ uses GraphQL (mobile tech from Facebook)

Aws TA (Trusted Advisor)

* Links: aws.amazon.com/architecture

" " " / solutions